

Advanced NLP has been completely transformed by massive AI models utilizing Transformer networks. The inherent self-attention process enables these systems to evaluate the semantic weight of individual terms within a phrase, irrespective of how far apart they are placed. In this study, we introduce a new method for streamlining the attention grid, which decreases processing requirements by forty percent while still preserving a ninety-five percent success rate on typical evaluation metrics such as SuperGLUE and GLUE. Our approach substitutes regular dense matrices with sparse data structures throughout the initial training period.